

demo_notebook

June 3, 2021

Colab <-> Github Credit: <https://medium.com/@ashwindesilva/how-to-use-google-colaboratory-to-clone-a-github-repository-e07cf8d3d22b>

<https://github.com/jakee417/ProbabilisticForecasting> for our full repo

```
[1]: from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
[3]: % cd gdrive/My Drive/Stat 271/
```

/content/gdrive/My Drive/Stat 271

```
[4]: ! git clone https://github.com/jakee417/ProbabilisticForecasting
```

```
Cloning into 'ProbabilisticForecasting'...
remote: Enumerating objects: 393, done.
remote: Counting objects: 100% (393/393), done.
remote: Compressing objects: 100% (345/345), done.
remote: Total 393 (delta 47), reused 390 (delta 47), pack-reused 0
Receiving objects: 100% (393/393), 61.22 MiB | 13.94 MiB/s, done.
Resolving deltas: 100% (47/47), done.
Checking out files: 100% (377/377), done.
```

```
[5]: % cd ProbabilisticForecasting
```

/content/gdrive/My Drive/Stat 271/ProbabilisticForecasting

```
[6]: from WindowGenerator import WindowGenerator
import matplotlib.pyplot as plt
from LstmRnn import LstmRnn
import json
import time
import pprint
from datetime import datetime
pp = pprint.PrettyPrinter(indent=4)
```

```
[7]: # forecast example:
# input_width=96,
# label_width=24,
# shift=24,
```

```
# anomaly example:
# input_width=24,
# label_width=24,
# shift=0,
```

```
[9]: # algorithm parameters
bitcoin = 'data/bitcoin_query.csv'
ethereum = 'data/ethereum_clean.csv'
# distributions: 'normal', 'student_t', 'mix', 'laplace', 'poisson_approx'
params = dict(
    fname=ethereum,
    distribution='normal', # likelihood model
    lstm_units=32, # hidden units for lstm and dense cells. #was 32
    t2v_units=8, # might leave this to 8 for now.
    dense_cells=None, # extra dense cells after lstm cell. Can be None.
    resample=None, # leave this none for more data
    input_width=24, # past 3 days
    label_width=24, # one day
    shift=0, # determines the index of the first forecast point . 0 for
    ↪ anomaly, label_width for forecast.
    max_epochs=40,
    patience=2, # reduce_on_plateau=patience, early_stopping=patience*2
    latent_dim=None, # can be None or 2 and above. #was 2
    beta=1, # 1 or greater
    min_df=2.0, # for the student-t
    number_states=30, # for the hmm
    batch_size=256, # 256 => 30s/epoch
    regularization= None # applies to lstm and dense cell kernels. Can be
    ↪ None.
)
```

```
[10]: now = datetime.now().strftime("%d_%m_%Y_%H_%M_%S")
distribution = params['distribution']
checkpoint_path = f'checkpoints/{distribution}.ckpt'
save_path = f'saved/LstmRnn{distribution}'
train_save_img = f'figures/{now}_{distribution}_train_lstm_rnn.jpg'
val_save_img = f'figures/{now}_{distribution}_val_lstm_rnn.jpg'
test_save_img = f'figures/{now}_{distribution}_test_lstm_rnn.jpg'
global_train_img = f'figures/{now}_train_{distribution}_global_lstm_rnn.jpg'
global_val_img = f'figures/{now}_val_{distribution}_global_lstm_rnn.jpg'
global_test_img = f'figures/{now}_test_{distribution}_global_lstm_rnn.jpg'
loss_img = f'figures/{now}_{distribution}_loss.jpg'
```

```

post_check_img = f'figures/{now}_{distribution}_post_check.jpg'
train_correlation_img = f'figures/{now}_train_{distribution}_correlation'
test_correlation_img = f'figures/{now}_test_{distribution}_correlation'

```

```

[11]: multi_window = WindowGenerator(params['fname'],
                                     input_width=params['input_width'],
                                     label_width=params['label_width'],
                                     shift=params['shift'],
                                     label_columns=['num_transactions'],
                                     resample_frequency=params['resample'],
                                     standardize=True,
                                     batch_size=params['batch_size'])

# multi_window.plot_splits(feedback_model)
print(multi_window)
time_index = multi_window.column_indices['timestamp']

```

```

Total window size: 24
Input indices: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
21 22 23]
Shift (Input) -> (Label): 0
Label indices: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
21 22 23]
Label column name(s): ['num_transactions']
Total time series: 38328

```

```

[12]: print(len(multi_window.train))
      print(len(multi_window.val))
      print(len(multi_window.test))

```

```


105
30
15

```

```

[13]: multi_window.plot_splits()

```



```
[14]: params.update({'time_index': time_index})
      feedback_model = LstmRnn(params)
      feedback_model(multi_window.train_example[0])
      print(feedback_model.summary())
```

Model: "lstm_rnn"

Layer (type)	Output Shape	Param #
lstm_warmup (LSTMCell)	multiple	5632
lstm_forecast (LSTMCell)	multiple	4480
rnn (RNN)	multiple	5632
t2v (T2V)	multiple	248
normal (DistributionLambda)	multiple	0
dense_extra (Dense)	multiple	0 (unused)
dense_main (Dense)	multiple	66

Total params: 10,434

Trainable params: 10,426
Non-trainable params: 8

None

```
[15]: tic = time.time()
      history = feedback_model.compile_and_fit(
          model=feedback_model,
          window=multi_window,
          checkpoint_path=checkpoint_path,
          save_path=save_path,
          max_epochs=params['max_epochs'],
          patience=params['patience']
      )
      toc = time.time()
```

Epoch 1/40

105/105 [=====] - 14s 74ms/step - loss: 32.4046 -
reconstruction_loss: 26.6951 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 37.5125 - val_reconstruction_loss: 37.5125 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00001: saving model to checkpoints/normal.ckpt

Epoch 2/40

105/105 [=====] - 7s 69ms/step - loss: 17.9934 -
reconstruction_loss: 17.2515 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 19.6801 - val_reconstruction_loss: 19.6801 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00002: saving model to checkpoints/normal.ckpt

Epoch 3/40

105/105 [=====] - 7s 66ms/step - loss: 15.5020 -
reconstruction_loss: 15.1190 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 17.4951 - val_reconstruction_loss: 17.4951 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00003: saving model to checkpoints/normal.ckpt

Epoch 4/40

105/105 [=====] - 7s 66ms/step - loss: 13.8948 -
reconstruction_loss: 13.2652 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 15.4511 - val_reconstruction_loss: 15.4511 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00004: saving model to checkpoints/normal.ckpt

Epoch 5/40

105/105 [=====] - 7s 67ms/step - loss: 11.0062 -
reconstruction_loss: 10.1933 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 13.9578 - val_reconstruction_loss: 13.9578 -

val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00005: saving model to checkpoints/normal.ckpt

Epoch 6/40

105/105 [=====] - 7s 66ms/step - loss: 7.9875 -
reconstruction_loss: 7.4063 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 13.8946 - val_reconstruction_loss: 13.8946 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00006: saving model to checkpoints/normal.ckpt

Epoch 7/40

105/105 [=====] - 7s 64ms/step - loss: 5.8498 -
reconstruction_loss: 5.4163 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 11.9167 - val_reconstruction_loss: 11.9167 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00007: saving model to checkpoints/normal.ckpt

Epoch 8/40

105/105 [=====] - 7s 65ms/step - loss: 4.2411 -
reconstruction_loss: 3.8753 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 9.0213 - val_reconstruction_loss: 9.0213 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00008: saving model to checkpoints/normal.ckpt

Epoch 9/40

105/105 [=====] - 7s 63ms/step - loss: 2.8876 -
reconstruction_loss: 2.6351 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 7.6950 - val_reconstruction_loss: 7.6950 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00009: saving model to checkpoints/normal.ckpt

Epoch 10/40

105/105 [=====] - 7s 63ms/step - loss: 1.8763 -
reconstruction_loss: 1.6688 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 7.0575 - val_reconstruction_loss: 7.0575 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00010: saving model to checkpoints/normal.ckpt

Epoch 11/40

105/105 [=====] - 7s 63ms/step - loss: 1.0558 -
reconstruction_loss: 0.8293 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 6.6502 - val_reconstruction_loss: 6.6502 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00011: saving model to checkpoints/normal.ckpt

Epoch 12/40

105/105 [=====] - 7s 61ms/step - loss: 0.3132 -
reconstruction_loss: 0.1916 - regularization_loss: 0.0000e+00 - kl_loss:

0.0000e+00 - val_loss: 5.7913 - val_reconstruction_loss: 5.7913 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00012: saving model to checkpoints/normal.ckpt

Epoch 13/40

105/105 [=====] - 7s 66ms/step - loss: -0.1923 -
reconstruction_loss: -0.3423 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 5.2229 - val_reconstruction_loss: 5.2229 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00013: saving model to checkpoints/normal.ckpt

Epoch 14/40

105/105 [=====] - 7s 62ms/step - loss: -0.5982 -
reconstruction_loss: -0.7891 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 5.6969 - val_reconstruction_loss: 5.6969 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00014: saving model to checkpoints/normal.ckpt

Epoch 15/40

105/105 [=====] - 7s 64ms/step - loss: -1.0749 -
reconstruction_loss: -1.1928 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 6.0275 - val_reconstruction_loss: 6.0275 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00015: ReduceLROnPlateau reducing learning rate to 0.00010000000474974513.

Epoch 00015: saving model to checkpoints/normal.ckpt

Epoch 16/40

105/105 [=====] - 7s 62ms/step - loss: -1.5423 -
reconstruction_loss: -1.5719 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 6.1370 - val_reconstruction_loss: 6.1370 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00

Epoch 00016: saving model to checkpoints/normal.ckpt

Epoch 17/40

105/105 [=====] - 7s 63ms/step - loss: -1.5712 -
reconstruction_loss: -1.6138 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00 - val_loss: 6.3181 - val_reconstruction_loss: 6.3181 -
val_regularization_loss: 0.0000e+00 - val_kl_loss: 0.0000e+00


Restoring model weights from the end of the best epoch.

Epoch 00017: ReduceLROnPlateau reducing learning rate to 1.0000000474974514e-05.

Epoch 00017: saving model to checkpoints/normal.ckpt


Epoch 00017: early stopping

```
[16]: # plot history of loss and val_loss
plt.figure(figsize=(12, 8))
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.legend(['loss', 'val_loss'])
plt.xlabel('Epochs')
plt.savefig(loss_img)
plt.show()
```




```
[17]: samples = 500
```

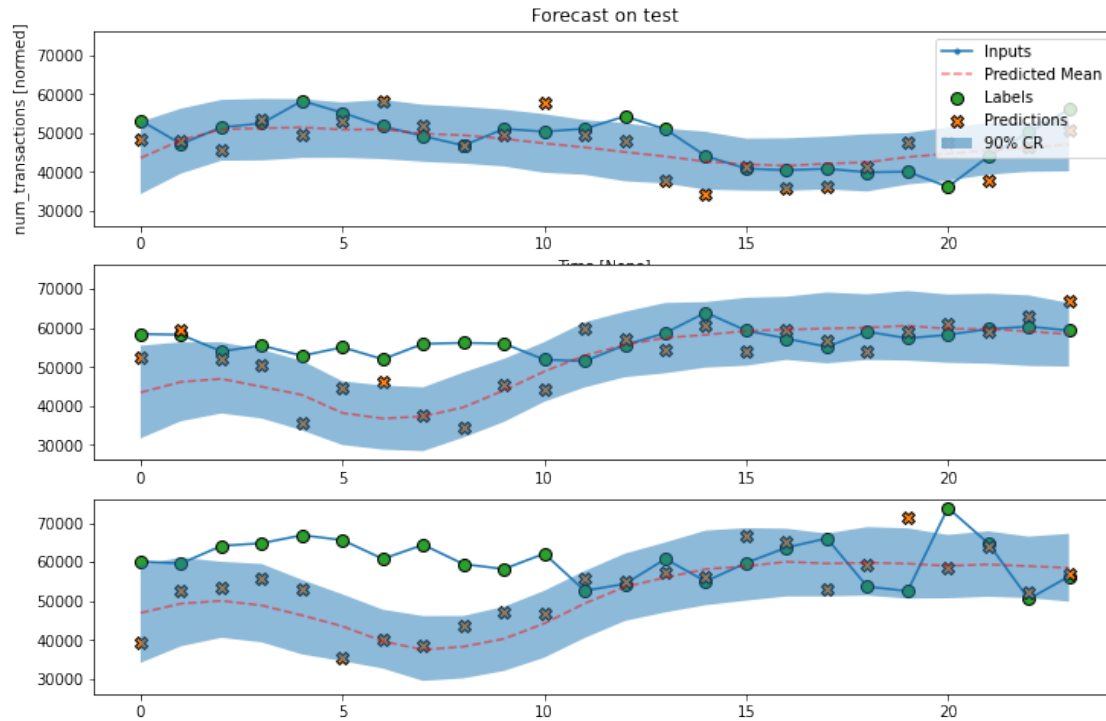
```
[18]: multi_window.plot(feedback_model,
                        save_path=train_save_img,
                        max_subplots=3,
                        samples=samples,
                        mode='train')
```

```
[19]: multi_window.plot(feedback_model,
                        save_path=val_save_img,
                        max_subplots=3,
                        samples=samples,
                        mode='val')
```



```
[20]: multi_window.plot(feedback_model,
                        save_path=test_save_img,
                        max_subplots=3,
                        samples=samples,
                        mode='test')
```



```
[21]: forecasts = []
```

```
[22]: # This takes a while - also don't need for now
train_forecast = multi_window.forecast(
    model=feedback_model,
    dataset_name='train',
    samples=samples
)
forecasts.append(train_forecast)
```

Sampling forecast values: train
 Sampling complete: 20.49944806098938
 Assign samples time index: 103.4306411743164

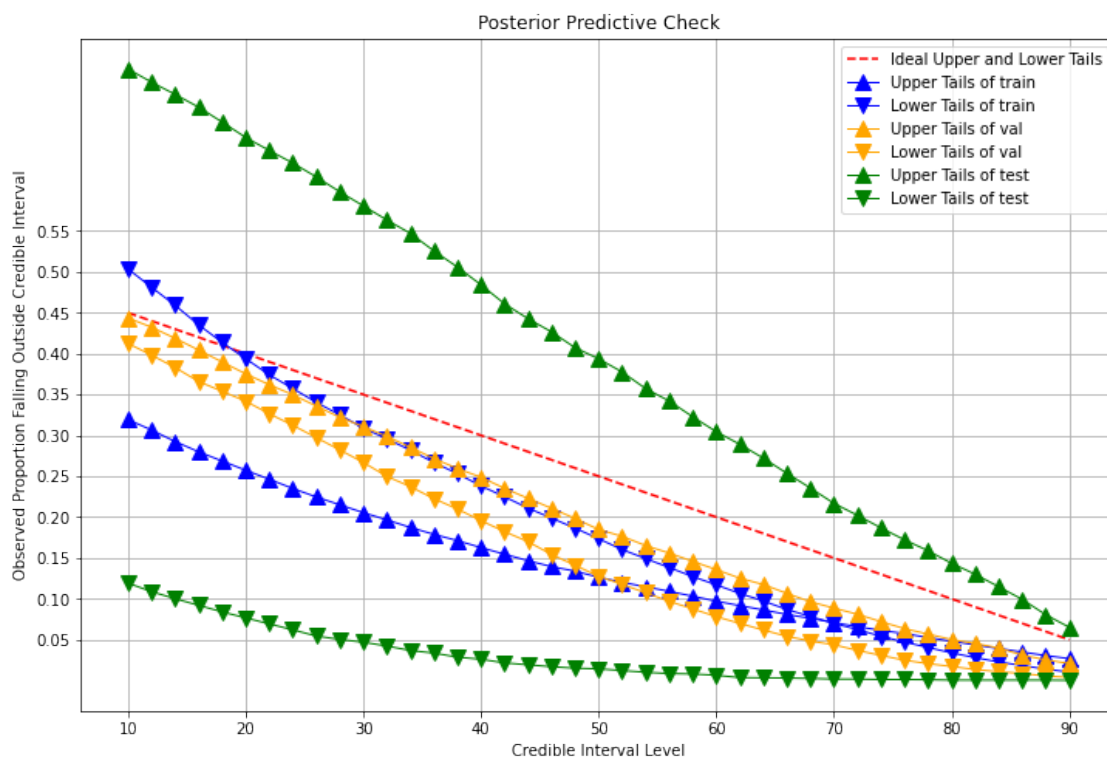
```
[23]: val_forecast = multi_window.forecast(
    model=feedback_model,
    dataset_name='val',
    samples=samples
)
forecasts.append(val_forecast)
```

Sampling forecast values: val
 Sampling complete: 5.8086793422698975
 Assign samples time index: 18.315253257751465


```
[24]: test_forecast = multi_window.forecast(
        model=feedback_model,
        dataset_name='test',
        samples=samples
    )
    forecasts.append(test_forecast)
```

Sampling forecast values: test
 Sampling complete: 5.136912822723389
 Assign samples time index: 10.379784345626831


```
[25]: post_checks = multi_window.plot_posterior_predictive_check(forecasts,
    ↪ post_check_img)
```




```
[26]: multi_window.plot_global_forecast(
        train_forecast,
        save_path=global_train_img
    )
```



```
[27]: multi_window.plot_global_forecast(  
    val_forecast,  
    save_path=global_val_img  
)
```



```
[28]: multi_window.plot_global_forecast(  
    test_forecast,  
    save_path=global_test_img  
)
```



```
[29]: # val_total_area, val_loss, test_total_area, and test_loss are the ones we care
      ↪ about for results
performance = dict()
performance['time'] = toc - tic
performance['now'] = now
performance.update(params)
performance['train_size'] = len(multi_window.train)
performance['val_size'] = len(multi_window.val)
performance['test_size'] = len(multi_window.test)
performance['total_data'] = len(multi_window.df)
performance['train'] = feedback_model.evaluate(multi_window.train)
performance['val'] = feedback_model.evaluate(multi_window.val)
performance['test'] = feedback_model.evaluate(multi_window.test, verbose=0)
performance.update(post_checks)
pp.pprint(performance)

# cache performance - saves automatically organized by datetime, distribution
out_file = open(f'metrics/{now}_{distribution}.json', "w")
json.dump(performance, out_file, indent=6)
out_file.close()
```

```
105/105 [=====] - 4s 35ms/step - loss: -0.5770 -
reconstruction_loss: -0.6197 - regularization_loss: 0.0000e+00 - kl_loss:
0.0000e+00
```

```
30/30 [=====] - 1s 36ms/step - loss: 5.2584 -  
reconstruction_loss: 5.2157 - regularization_loss: 0.0000e+00 - kl_loss:  
0.0000e+00
```

```
{  'batch_size': 256,  
    'beta': 1,  
    'dense_cells': None,  
    'distribution': 'normal',  
    'fname': 'data/ethereum_clean.csv',  
    'input_width': 24,  
    'label_width': 24,  
    'latent_dim': None,  
    'lstm_units': 32,  
    'max_epochs': 40,  
    'min_df': 2.0,  
    'now': '02_06_2021_03_07_35',  
    'number_states': 30,  
    'patience': 2,  
    'regularization': None,  
    'resample': None,  
    'shift': 0,  
    't2v_units': 8,  
    'test': [31.77865982055664, 0.0, 31.77865982055664, 0.0],  
    'test_area_above': 6.155948343334203,  
    'test_area_below': 9.062157578919907,  
    'test_size': 15,  
    'test_total_area': 15.22,  
    'time': 140.401602268219,  
    'time_index': 1,  
    'total_data': 38328,  
    'train': [-0.6197210550308228, 0.0, -0.6197210550308228, 0.0],  
    'train_area_above': 4.389699578813969,  
    'train_area_below': 2.3493067203399307,  
    'train_size': 105,  
    'train_total_area': 6.74,  
    'val': [5.215722560882568, 0.0, 5.215722560882568, 0.0],  
    'val_area_above': 1.8889251239238196,  
    'val_area_below': 3.6803417688494653,  
    'val_size': 30,  
    'val_total_area': 5.57}
```

```
[ ]:
```